

# Applied Data Science - Foundations module - Final Project

by Peter Varshavsky

## The data

In this project I analyze hourly bikeshare data from Capital Bikeshare in Washington, DC. The dataset includes hourly counts of bike use by subscribers and one-time users, as well as quantitative and categorical weather variables.

### *Variables*

- season:
  1. spring
  2. summer
  3. fall
  4. winter
- holiday: boolean
- workingday: boolean
- weather
  1. Clear, Few clouds, Partly cloudy, Partly cloudy
  2. Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
  3. Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
  4. Heavy Rain + Ice Pellets + Thunderstorm + Mist, Snow + Fog
- temp: degrees C
- atemp: 'feels like', degrees C
- humidity: relative humidity
- windspeed: wind speed
- casual: number of non-registered user rentals initiated
- registered: number of registered user rentals initiated
- count\_tot: number of total rentals

### *Training and test sets*

The data has been split into the training set containing hourly data for 1st-19th of each months, and test set containing data for the remainder of each month.

## Results

I modeled hourly variation and daily variation separately assuming they were multiplicatively related. I computed hourly averages separately for working days and non working days, as there is a clear difference in use patterns. Daily variation was modeled with a linear model using several subsets of variables, with the final model including *day\_num*, *season*, *temp*, *windspeed* *weather*.

The final model was submitted to Kaggle and received a Root Mean Squared Logarithmic Error score of 0.60530, ranking 1000th out of roughly 1500 participants. Because the test values used for scoring are not provided, it is not possible to visualize the residuals.

998	↓61	Chris H	0.60472	3	Mon, 13 Oct 2014 23:25:44
999	↓61	asda	0.60505	4	Wed, 26 Nov 2014 23:01:48 (-0.3h)
1000	new	<b>Peter Varshavsky</b>	<b>0.60530</b>	<b>1</b>	<b>Tue, 09 Dec 2014 21:48:31</b>
1001	↓62	Lotte	0.60536	3	Tue, 07 Oct 2014 14:30:19 (-11d)
1002	↓62	Simon Parker	0.60619	11	Thu, 09 Oct 2014 10:54:09

## Processing and modeling

I chose to fill the missing days of every month with the average of five days before and five days after each missing period. This allowed me to work with an uninterrupted signal which made it easier to attempt harmonic analysis with FFT. I briefly attempted using FFT in Python and ARIMA in R, but, due to lack of experience, decided to fit the periodic components using least squares with available variables (*season*, *datetime*).

Filling in the missing data may also lead to unreliable P-values and  $R^2$ , since the model thinks that it has more degrees of freedom than there truly are.

## Ideas for improvement of this approach

The model, as it is implemented has a number of problems. My next step would be refitting the model without interpolating the missing data, and inspecting the strangely outlying predicted counts. For example the counts predicted for January 21, 2011 are negative. For the submission I substituted all negative counts with zeros.

The weather is only used to fit daily use in this model, so the model is losing a lot of detail because hourly changes in temperature, wind, and precipitation are likely to have a lot of explanatory power.

## Other approaches

After fixing the obvious problems with the current implementation, I would want to learn how to use the customary time series tools like ARMA and ARIMA models.

In [1]:

```
import patsy
import os
import sys
import csv
```

```

import timeit
from datetime import datetime
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import calendar

%matplotlib inline
plt.style.use('ggplot')

```

```

In [2]: def readData_noParse(fileIn):
        train = pd.read_csv(fileIn, index_col = ['datetime'])
        train.index = pd.DatetimeIndex(train.index)
        return train

fileIn = "../data/train.csv"

train = readData_noParse(fileIn)
columns = list(train.columns)
columns[-1] = 'count_tot'
train.columns = columns

### Season recording is better after resample
# seasonsDict = {1: 'spring', 2: 'summer', 3: 'fall', 4: 'winter'}
# train['season'] = train.season.apply(lambda s: seasonsDict[s])
train.head(3)

```

```

Out[2]:

```

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	ca
2011-01-01 00:00:00	1	0	0	1	9.84	14.395	81	0	3
2011-01-01 01:00:00	1	0	0	1	9.02	13.635	80	0	8
2011-01-01 02:00:00	1	0	0	1	9.02	13.635	80	0	5

### Resample daily

```

In [3]: train_daily = pd.DataFrame(train.resample('D', how = 'mean'))
        print "Days without data (NaN):", sum(train_daily.count_tot.isnull())

Days without data (NaN): 263

```

### Missing days:

Fill in missing days (20th - EOM of each month) with an average of 5 days before and 5 days after.

```
In [4]: #print pd.concat([pd.DataFrame(train_daily.loc['2011-01-10':'2011-02-10'].mean()).transpose()* 11, ignore_index = True)
def fill_NA_days(df, sampleStart, sampleEnd, startNA, endNA):
#     average_list = list(train_daily.loc['2011-01-15':'2011-02-5'].mean())
#     average_list_of_lists = [average_list for x in range(12)]
#     train_daily.loc['2011-01-20':'2011-01-31'] = average_list_of_lists
    ndays = int(endNA[8:]) - 19
    average_list = list(df.loc[sampleStart:sampleEnd].mean())
    average_list_of_lists = [average_list for x in range(ndays)]
    df.loc[startNA:endNA] = average_list_of_lists
```

```
In [5]: months = [str(month).zfill(2) for month in range(1,13)] * 2
years = ['2011'] * 12 + ['2012'] * 12
sampleStarts = [years[i] + '-' + months[i] + '-' + '15' for i in range(len(months))]
sampleEnds = [years[i] + '-' + months[i] + '-' + '05' for i in range(1, len(months))]
startNAs = [years[i] + '-' + months[i] + '-' + '20' for i in range(len(months))]
monthEnds = [str(calendar.monthrange(2011, month)[1]).zfill(2) for month in range(1, 13)] + [str(calendar.monthrange(2012, month)[1]).zfill(2) for month in range(1, 13)]
endNAs = [years[i] + '-' + months[i] + '-' + monthEnds[i] for i in range(len(months))]

for i in range(len(sampleStarts) - 1):
    fill_NA_days(train_daily, sampleStarts[i], sampleEnds[i], startNAs[i], endNAs[i])

print "Number of missing values after filling in: %d" %sum(train_daily.count_tot.isnull())
```

Number of missing values after filling in: 0

```
In [6]: ### Note: can be rewritten simpler with Timestamp.month and wrapped in to a function
### to be used later on test set as well
### Warnings can be ignored
train_daily.season['2011-01':'2011-02'] = 'winter'
train_daily.season['2011-12':'2012-02'] = 'winter'
train_daily.season['2011-03':'2011-05'] = 'spring'
train_daily.season['2012-03':'2012-05'] = 'spring'
train_daily.season['2011-06':'2011-08'] = 'summer'
train_daily.season['2012-06':'2012-08'] = 'summer'
train_daily.season['2011-09':'2011-11'] = 'fall'
train_daily.season['2012-09':'2012-11'] = 'fall'
train_daily.season['2012-12'] = 'winter'
```

-c:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:7: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:8: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:9: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:10: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:11: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:12: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

## Linear model with 1 variable *day\_num*

This model explains a fair amount of daily variation with  $R^2 = 0.563$ , but it's obvious from the plots of the model and the residuals, that it can be improved with a seasonal component.

```
In [7]: train_daily['day_num'] = list(range(1, train_daily.shape[0] + 1))
        print train_daily.head(3)
        print train_daily.columns

y, X = patsy.dmatrices('count_tot ~ day_num', data = train_daily)
mod1 = sm.OLS(y, X).fit()
```

season holiday workingday weather temp atem

```

p \
2011-01-01  winter      0      0  1.583333  14.110833  18.18125
0
2011-01-02  winter      0      0  1.956522  14.902609  17.68695
7
2011-01-03  winter      0      1  1.000000   8.050909   9.47022
7

```

```

          humidity  windspeed      casual  registered  count_tot  da
y_num
2011-01-01  80.583333  10.749871  13.791667   27.250000  41.041667
1
2011-01-02  69.608696  16.652122   5.695652   29.130435  34.826087
2
2011-01-03  43.727273  16.636709   5.454545   55.863636  61.318182
3

```

```

Index([u'season', u'holiday', u'workingday', u'weather', u'temp', u'at
emp', u'humidity', u'windspeed', u'casual', u'registered', u'count_tot
', u'day_num'], dtype='object')

```

In [8]: `print mod1.summary()`

#### OLS Regression Results

```

=====
=====
Dep. Variable:          count_tot    R-squared:
    0.563
Model:                OLS    Adj. R-squared:
    0.562
Method:             Least Squares    F-statistic:
    923.3
Date:                Tue, 09 Dec 2014    Prob (F-statistic):
    .64e-131    5
Time:                19:18:26    Log-Likelihood:
    -3818.3
No. Observations:    719    AIC:
    7641.
Df Residuals:        717    BIC:
    7650.
Df Model:            1

```

```

=====
=====
              coef      std err          t      P>|t|      [95.0% Con
f. Int.]
-----
Intercept    96.3837      3.663      26.313      0.000      89.192
103.575
day_num      0.2679      0.009      30.386      0.000      0.251
0.285

```

```

=====
=====
Omnibus:          37.563    Durbin-Watson:

```

```

0.446
Prob(Omnibus):          0.000   Jarque-Bera (JB):
22.660
Skew:                  -0.291   Prob(JB):
1.20e-05
Kurtosis:              2.354   Cond. No.
832.
=====
=====

```

```

In [9]: fig1, [ax1_1, ax1_2] = plt.subplots(nrows = 2, figsize = (20, 8))
ax1_1.plot(train_daily.count_tot)
ax1_1.plot(mod1.predict())
ax1_2.plot(train_daily.count_tot - mod1.predict(), '.')
ax1_1.set_ylabel("Count")
ax1_2.set_ylabel("Residuals")
ax1_2.set_xlabel("Days")
ax1_1.set_xlim([0, 719])
ax1_2.set_xlim([0, 719])
fig1.suptitle("Figure 1. Model 1: linear function of time", fontsize =
14)
plt.show()

```

Figure 1. Model 1: linear function of time



## Model 2. Two variables: *day\_num*, *season*

Adding season as a categorical variable improves  $R^2$  to 0.784. As may be expected, there is more bike use in spring and summer, and less in winter than during fall (baseline).

```

In [10]: y, X = patsy.dmatrices('count_tot ~ day_num + season', data = train_da
ily)
mod2 = sm.OLS(y, X).fit()
print mod2.summary()

```

OLS Regression Results

```

=====
=====

```

```

Dep. Variable:          count_tot    R-squared:
0.784
Model:                  OLS          Adj. R-squared:
0.783
Method:                 Least Squares    F-statistic:
649.2
Date:                   Tue, 09 Dec 2014    Prob (F-statistic):      4
.06e-236
Time:                   19:18:26          Log-Likelihood:
-3564.3
No. Observations:      719              AIC:
7139.
Df Residuals:          714              BIC:
7162.
Df Model:               4

```

```

=====
=====

```

	coef	std err	t	P> t	[95.
Intercept	94.8055	4.051	23.406	0.000	86
season[T.spring]	24.7885	3.811	6.505	0.000	17
season[T.summer]	45.0108	3.661	12.294	0.000	37
season[T.winter]	-49.8802	3.872	-12.881	0.000	-57
day_num	0.2552	0.007	38.368	0.000	0

```

=====
=====

```

```

Omnibus:                35.300    Durbin-Watson:
0.928
Prob(Omnibus):          0.000    Jarque-Bera (JB):
44.729
Skew:                   -0.464    Prob(JB):
1.94e-10
Kurtosis:               3.794    Cond. No.
2.07e+03

```

```

=====
=====

```

```

Warnings:
[1] The condition number is large, 2.07e+03. This might indicate that
there are
strong multicollinearity or other numerical problems.

```

```

In [11]: fig2, [ax2_1, ax2_2] = plt.subplots(num = 2, nrows = 2, figsize = (20,
8))
ax2_1.plot(train_daily.count_tot)
ax2_1.plot(mod2.predict())

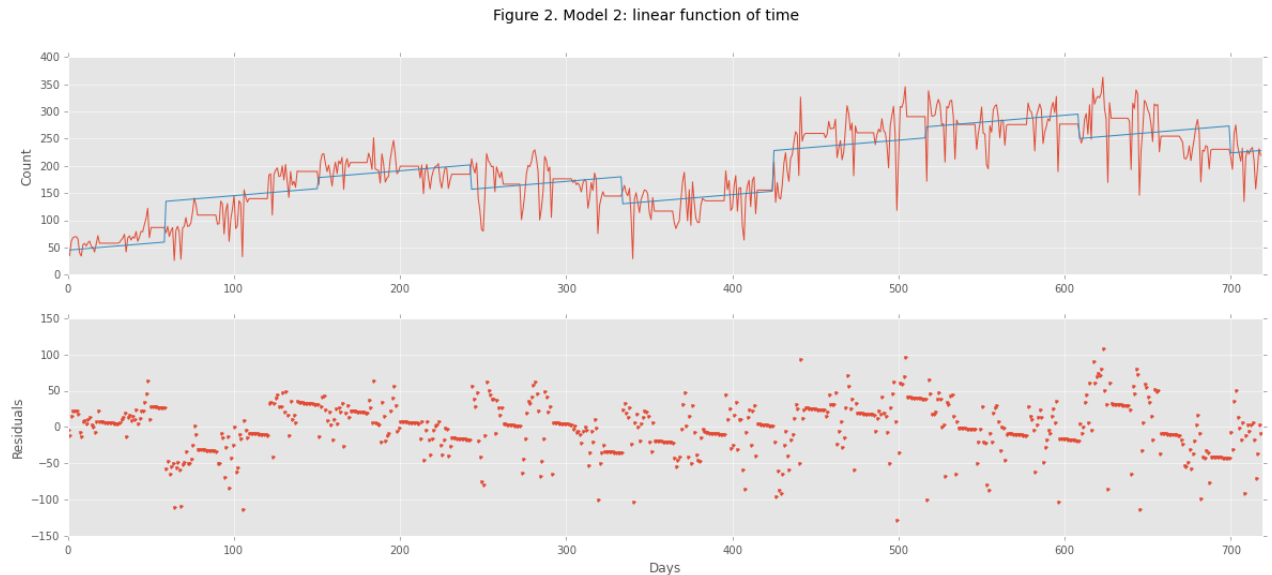
```



```

ax2_2.plot(train_daily.count_tot - mod2.predict(), '.')
ax2_1.set_ylabel("Count")
ax2_2.set_ylabel("Residuals")
ax2_2.set_xlabel("Days")
ax2_1.set_xlim([0, 719])
ax2_2.set_xlim([0, 719])
fig2.suptitle("Figure 2. Model 2: linear function of time", fontsize =
14)
plt.show()

```



## Model 3. Adding quantitative weather variables:

- Variables in model: *day\_num*, *season*, *temp*, *humidity*, *windspeed*

*atemp* is air temperature or 'feels like', which is highly correlated to *temp*, *humidity*, and *windspeed*. It is left out.

```

In [12]: y, X = patsy.dmatrices('count_tot ~ day_num + season + temp + humidity
+ windspeed', data = train_daily)
mod3 = sm.OLS(y, X).fit()
print mod3.summary()

```

### OLS Regression Results

```

=====
=====
Dep. Variable:          count_tot    R-squared:
    0.874
Model:                  OLS          Adj. R-squared:
    0.873
Method:                 Least Squares    F-statistic:
    705.5
Date:                   Tue, 09 Dec 2014    Prob (F-statistic):
    .28e-315
Time:                   19:18:27          Log-Likelihood:
    -3370.7
No. Observations:      719              AIC:

```

```

6757.
Df Residuals:          711    BIC:
6794.
Df Model:              7

```

```

=====
=====
                                coef      std err          t      P>|t|      [95.
0% Conf. Int.]
-----
Intercept          120.7840        9.299      12.989      0.000      102
.528    139.040
season[T.spring]    16.1397        3.003       5.375      0.000      10
.245    22.035
season[T.summer]   -15.9060        4.106      -3.874      0.000     -23
.967    -7.845
season[T.winter]   -18.2836        3.587      -5.097      0.000     -25
.326   -11.241
day_num            0.2258        0.005      42.907      0.000       0
.215     0.236
temp              5.2826        0.287      18.379      0.000       4
.718     5.847
humidity          -1.2782        0.094     -13.645      0.000      -1
.462    -1.094
windspeed         -2.6212        0.252     -10.414      0.000      -3
.115    -2.127
=====
=====
Omnibus:          141.592    Durbin-Watson:
1.181
Prob(Omnibus):    0.000    Jarque-Bera (JB):
668.910
Skew:            -0.808    Prob(JB):          5
.60e-146
Kurtosis:        7.440    Cond. No.
4.06e+03
=====
=====

```

#### Warnings:

```

[1] The condition number is large, 4.06e+03. This might indicate that
there are
strong multicollinearity or other numerical problems.


```

```

In [13]: fig3, [ax3_1, ax3_2] = plt.subplots(num = 3, nrows = 2, figsize = (20,
8))
ax3_1.plot(train_daily.count_tot)
ax3_1.plot(mod3.predict())
ax3_2.plot(train_daily.count_tot - mod3.predict(), '.')
ax3_1.set_ylabel("Count")
ax3_2.set_ylabel("Residuals")
ax3_2.set_xlabel("Days")
ax3_1.set_xlim([0, 719])
ax3_2.set_xlim([0, 719])

```

```
fig3.suptitle("Figure 3. Model 3: linear function of time", fontsize =
16)
plt.show()
```



## Model 4. Adding qualitative weather variable:

- Variables in model: *day\_num*, *season*, *temp*, *windspeed*, *weather*
- Variables taken out: *humidity*

With addition of *weather*, *humidity* lost significance.

**TODO:** Using ANOVA, test whether leaving out remaining variables: *humidity*, *workingday*, *holiday* improves the model.

```
In [14]: y, X = patsy.dmatrices('count_tot ~ day_num + season + temp + weather
+ windspeed', data = train_daily)
mod4 = sm.OLS(y, X).fit()
print mod4.summary()
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          count_tot    R-squared:
      0.891
Model:                  OLS          Adj. R-squared:
      0.890
Method:                 Least Squares    F-statistic:
      826.9
Date:                  Tue, 09 Dec 2014    Prob (F-statistic):
      0.00
Time:                  19:18:27          Log-Likelihood:
      -3320.3
No. Observations:      719              AIC:
      6657.
Df Residuals:          711              BIC:
      6693.
```

```
=====
=====
                                coef      std err          t      P>|t|      [95.
0% Conf. Int.]
-----
Intercept                106.6162        7.591      14.045      0.000        91
.713    121.520
season[T.spring]          21.5969        2.761       7.822      0.000        16
.176     27.018
season[T.summer]          -8.4998        3.689      -2.304      0.022       -15
.743     -1.256
season[T.winter]         -16.1235        3.341      -4.826      0.000       -22
.683     -9.564
day_num                   0.2315        0.005     47.492      0.000         0
.222      0.241
temp                     4.5303        0.261     17.368      0.000         4
.018      5.042
weather                  -47.9487        2.675    -17.922      0.000       -53
.201    -42.696
windspeed                -1.6305        0.224     -7.275      0.000        -2
.071     -1.190
=====
=====
Omnibus:                  37.925   Durbin-Watson:
    1.009
Prob(Omnibus):            0.000   Jarque-Bera (JB):
    107.341
Skew:                    -0.185   Prob(JB):
4.91e-24
Kurtosis:                4.856   Cond. No.
3.63e+03
=====
=====
```

## Warnings:

[1] The condition number is large, 3.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [15]: fig4, [ax4_1, ax4_2] = plt.subplots(num = 4, nrows = 2, figsize = (20,
      8))
ax4_1.plot(train_daily.count_tot)
ax4_1.plot(mod4.predict())
ax4_2.plot(train_daily.count_tot - mod4.predict(), '.')
ax4_1.set_ylabel("Count")
ax4_2.set_ylabel("Residuals")
ax4_2.set_xlabel("Days")
ax4_1.set_xlim([0, 719])
ax4_2.set_xlim([0, 719])
fig4.suptitle("Figure 4. Model 4: linear function of time", fontsize =
16)
plt.show()
```

Figure 4. Model 4: linear function of time



## Using *mod4* on test set

### Read test set

```
In [16]: fileIn = "../data/test.csv"

test = readData_noParse(fileIn)
test.index = pd.DatetimeIndex(test.index)
print "Test shape: ", test.shape
print test.head(3)

### Resample daily
test_daily = test.resample('D', how = 'mean')
print "test_daily shape (after resampling by date):", test_daily.shape
test_daily['day_num'] = (test_daily.index - pd.Timestamp('2011-01-01'))
    .days + 1
### Season names
test_daily.season['2011-01':'2011-02'] = 'winter'
test_daily.season['2011-12':'2012-02'] = 'winter'
test_daily.season['2011-03':'2011-05'] = 'spring'
test_daily.season['2012-03':'2012-05'] = 'spring'
test_daily.season['2011-06':'2011-08'] = 'summer'
test_daily.season['2012-06':'2012-08'] = 'summer'
test_daily.season['2011-09':'2011-11'] = 'fall'
test_daily.season['2012-09':'2012-11'] = 'fall'
test_daily.season['2012-12'] = 'winter'
### Drop Nan
test_daily.dropna(inplace = True)
print "test_daily shape after dropping NaNs: ", test_daily.shape
```

-c:13: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:14: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:15: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:16: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:17: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:18: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:20: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

-c:21: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Test shape: (6493, 8)

	season	holiday	workingday	weather	temp	ate
mp \						
2011-01-20 00:00:00	1	0	1	1	10.66	11.3
65						
2011-01-20 01:00:00	1	0	1	1	10.66	13.6
35						
2011-01-20 02:00:00	1	0	1	1	10.66	13.6
35						

	humidity	windspeed
2011-01-20 00:00:00	56	26.0027
2011-01-20 01:00:00	56	0.0000
2011-01-20 02:00:00	56	0.0000

test\_daily shape (after resampling by date): (712, 8)

test daily shape after dropping NaNs: (275, 9)

## Predicting test set daily totals

```
In [17]: _, X = patsy.dmatrices('day_num ~ day_num + season + temp + weather +
windspeed', data = test_daily)
test_daily['mod4'] = (mod4.predict(X) * 24).astype(int)
```

## Within-day variation

### TODO:

- Currently averages are divided by 24. This doesn't make sense

```
In [28]: seasonsDict = {1: 'spring', 2: 'summer', 3: 'fall', 4: 'winter'}

hourlyBySW = train.groupby(['workingday', 'season', lambda x: x.hour])
            .count_tot.mean()

hourlyByW = train.groupby(['workingday', lambda x: x.hour]).mean()

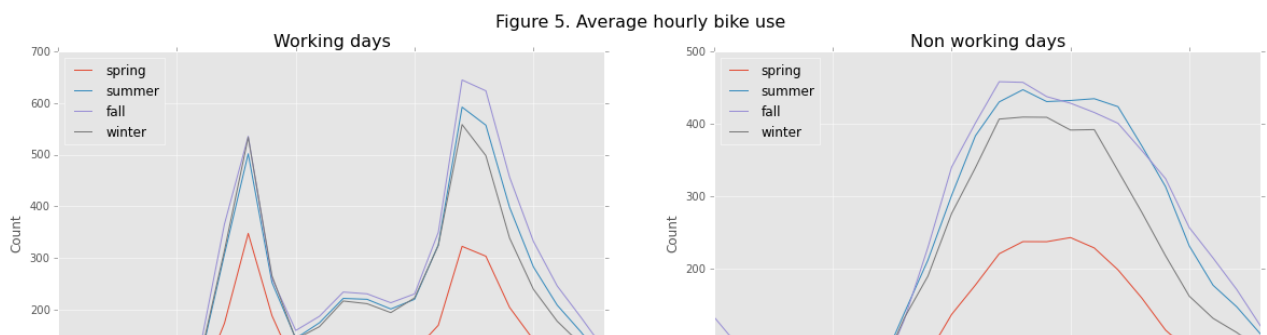
fig5, axs5 = plt.subplots(ncols = 2, num = 5, figsize = (20,6))



for season in [1,2,3,4]:
    ### workingday == 1
    axs5[0].plot(hourlyBySW.loc[1,season,:], label = seasonsDict[season])
    ### workingday == 0
    axs5[1].plot(hourlyBySW.loc[0,season,:], label = seasonsDict[season])

for ax in axs5:
    ax.legend(loc = 2)
    ax.set_xlabel('Hour')
    ax.set_xlim((0,23))
    ax.set_ylabel('Count')

axs5[0].set_title('Working days', fontsize = 16)
axs5[1].set_title('Non working days', fontsize = 16)
fig5.suptitle("Figure 5. Average hourly bike use", fontsize = 16)

plt.show()
```





# Making hourly predictions

## Computing average fractions of daily use for each hour

```
In [29]: ### computing average hourly use
averageDailySum = [sum(hourlyByW.loc[workday,:].count_tot) for workday
    in [0,1]]
hourlyFractions = [hourlyByW.loc[workday,:].count_tot / averageDailySum[workday] for workday in [0, 1]]
hourlyFractionsDict = {workday:{i:fraction for i, fraction in hourlyFractions[workday].iteritems()} for workday in [0,1]}
### plotting average hourly use
fig6, axs6 = plt.subplots(num = 6, ncols = 2, figsize = (20,6))

axs6[0].plot(hourlyFractions[1])
axs6[1].plot(hourlyFractions[0])


for ax in axs6:
    ax.set_xlim((0,23))

axs6[0].set_title('Working days', fontsize = 16)
axs6[1].set_title('Non working days', fontsize = 16)

for ax in axs6:
    ax.set_xlabel('Hour')
    ax.set_xlim((0,23))

fig6.suptitle("Figure 6. Normalized Distribution of average hourly bike use", fontsize = 16)

plt.show()
```



## Resampling back to hourly resolution

- Resample mod4 predictions by hour and fill NA with daily counts



- Note: mod4\_hourly picks up all dates, not just the 20th-EOM of each month these dates will be eliminated in joining with original test file

```
In [30]: mod4_hourly = test_daily[['mod4']].resample('H', how = 'sum').fillna(method = 'ffill')
### Join mod4_hourly with raw test data
#print "test rows: %d\nmod4 rows: %d" %(test.shape[0], mod4_hourly.shape[0])
test_prediction = pd.merge(test, mod4_hourly, how = 'inner', left_index = True, right_index = True)
#print "Merged shape: %d" %test_prediction.shape[0]
### Hour column
test_prediction['hour'] = test_prediction.index.hour
test_prediction['workdayPred'] = test_prediction.hour.replace(hourlyFractionsDict[1], inplace=False) * test_prediction.workingday * test_prediction.mod4
test_prediction['nonworkdayPred'] = test_prediction.hour.replace(hourlyFractionsDict[0], inplace=False) * (1 - test_prediction.workingday) * test_prediction.mod4
test_prediction['hourlyPred'] = test_prediction.workdayPred + test_prediction.nonworkdayPred
test_prediction.head(3)
```

Out[30]:

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	m
2011-01-20 00:00:00	1	0	1	1	10.66	11.365	56	26.0027	11
2011-01-20 01:00:00	1	0	1	1	10.66	13.635	56	0.0000	11
2011-01-20 02:00:00	1	0	1	1	10.66	13.635	56	0.0000	11

## Cleanup of submission

- Setting negative values to 0 (need to investigate negative values)
- Loading sample submissions into a data frame and joining with predictions

```
In [31]: ### Final submission series
final_predictions = pd.DataFrame(test_prediction.hourlyPred.astype(int))
final_predictions.columns = ['count']
final_predictions[final_predictions < 0] = 0

### Load sample file
fname = "sampleSubmission.csv"
fIn = os.path.join("../", "data", fname)
sample = pd.read_csv(fIn, parse_dates = [0])
sample.index = sample.datetime
```

```

sample.columns
sample.drop('datetime', 1, inplace = True)
final_out = pd.merge(final_predictions, sample, how = 'right', left_in
dex = True, right_index = True)
final_out.drop(u'count_y', 1, inplace = True)
final_out = pd.Series(final_out.count_x)
final_out.name = 'count'
final_out[final_out.isnull()] = 0

```

### ***Write hourly predictions to file***

```

In [32]: outputFolder = "output"
time = datetime.strftime(datetime.now(), "%Y-%m-%d_%H-%M-%S")
outputFileName = "prediction_file_" + time + ".csv"
fOut = os.path.join("../", outputFolder, outputFileName)
final_out.to_csv(fOut, header = True, index_label = "datetime")

```

### **Quick fixes to February 26, 2011 (negative values) and last week of Dec, 2012 (not in my predictions due to a coding error)**

To do.

```

In [33]: twoPrior = np.array(final_out['2011-1-24':'2011-1-25'].resample('H', h
ow = 'mean'))
twoAfter = np.array(final_out['2011-1-27':'2011-1-28'].resample('H', h
ow = 'mean'))

# print 'a', twoPrior.shape
# print 'a', twoAfter.shape

quickFix26 = np.zeros(24)
# for date in ['24', '25', '26', '27', '28']:
#     #quickFix26 += np.array(final_out['2011-01-'+date])
#     print final_out['2011-1-' + date].shape
final_out['2011-1-26']

```

```

Out[33]: datetime
2011-01-26 00:00:00    0
2011-01-26 01:00:00    0
2011-01-26 02:00:00    0
2011-01-26 05:00:00    0
2011-01-26 06:00:00    0
2011-01-26 07:00:00    0
2011-01-26 08:00:00    0
2011-01-26 09:00:00    0
2011-01-26 10:00:00    0
2011-01-26 11:00:00    0
2011-01-26 12:00:00    0
2011-01-26 13:00:00    0
2011-01-26 14:00:00    0
2011-01-26 15:00:00    0
2011-01-26 16:00:00    0
2011-01-26 17:00:00    0

```

Name: count, dtype: float64

In [ ]: